

# What is Data Retrieval?



**SDSC** SAN DIEGO  
SUPERCOMPUTER CENTER

# What is Data Retrieval?

- **Data retrieval**

- The way in which the desired data is specified and retrieved from a data store

- **Our focus**

- How to specify a data request
  - For static and streaming data
- The internal mechanism of data retrieval
  - For large and streaming data

# What is a Query Language?

- A language to specify the data items you need
- **A query language is declarative**
  - Specify what you need rather than how to obtain it
  - SQL (Structured Query Language)
- **Database programming language**
  - Procedural programming language
  - Embeds query operations

# SQL

- The standard for structured data
  - Oracle's SQL to Spark SQL
- Example Database Schema

**Bars(name, addr, license)**

**Beers(name, manf)**

**Sells(bar, beer, price)**

**Drinkers(name, addr, phone)**

**Frequents(drinker, bar)**

**Likes(drinker, beer)**

<u>name</u>	<u>addr</u>	<u>license</u>
Great American Bar	363 Main St., SD, CA 92390	41-437844098
Beer Paradise	6450 Mango Drive, SD, CA 92130	41-973428319
Have a Good Time	8236 Adams Avenue, SD, CA 92116	32-032263401

# SELECT-FROM-WHERE

- Which beers are made by Heineken?

```
SELECT name  
FROM Beers  
WHERE manf = 'Heineken'
```

*Output attribute(s)*

*Table(s) to use*

*The condition(s) to satisfy*

*Strings like 'Heineken' are case-sensitive and are put in quotes*

name
Heineken Lager Beer
Amstel Lager
Amstel Light
...

*Select<sub>manf='Heineken'</sub> (Beers)*



*Project(name)*

# More Example Queries

- Find expensive beer
  - SELECT DISTINCT beer, price
  - FROM Sells
  - WHERE price > 15
- Which businesses have a Temporary License (starts with 32) in San Diego?
  - SELECT name
  - FROM Bars
  - WHERE addr LIKE '%SD%' **AND** license LIKE '32%' LIMIT 5

<u>name</u>	<u>addr</u>	<u>license</u>
Great American Bar	363 Main St., SD, CA 92390	41-437844098
Beer Paradise	6450 Mango Drive, SD, CA 92130	41-973428319
Have a Good Time	8236 Adams Avenue, SD, CA 92116	32-032263401

# Select-Project Queries in the Large

- Large Tables can be partitioned
  - Many partitioning schemes
    - Range partitioning on primary key

name	manf	name	manf	...	name	manf	...
A...	Gambrinus	C...	MillerCoors		H...	Heineken	
A...	Heineken	C...	MillerCoors		H...	Pabst	
...		...			...		
B...	Anheuser-Busch	D...	Duvel Moortgat		H...	Anheuser-Busch	
Machine 1		Machine 2			Machine 5		

# Select-Project Queries in the Large

name	manf	name	manf	...	name	manf
A...	Gambrinus	C...	MillerCoors		H...	Heineken
A...	Heineken	C...	MillerCoors		H...	Pabst
...		...			...	
B...	Anheuser-Busch	D...	Duvel Moortgat		H...	Anheuser-Busch
Machine 1		Machine 2			Machine 5	

```
SELECT *  
FROM Beers  
WHERE name like 'Am%'
```

*pattern*

```
SELECT name  
FROM Beers  
WHERE manf = 'Heineken'
```

- **Two queries**

- Find records for beers whose name starts with 'Am'
- Which beers are made by Heineken?



# Evaluating SP Queries for Large Data

name	manf	name	manf	...	name	manf
A...	Gambrinus	C...	MillerCoors		H...	Heineken
A...	Heineken	C...	MillerCoors		H...	Pabst
...		...			...	
B...	Anheuser-Busch	D...	Duvel Moortgat		H...	Anheuser-Busch
Machine 1		Machine 2			Machine 5	

```
SELECT *  
FROM Beers  
WHERE name like 'Am%'
```

- A query processing trick
  - Use the partitioning information
    - Just use partition 1!!

# Evaluating SP Queries for Large Data

name	manf	name	manf	name	manf
A...	Gambrinus	C...	MillerCoors	H...	Heineken
A...	Heineken	C...	MillerCoors	H...	Pabst
...		...		...	
B...	Anheuser-Busch	D...	Duvel Moortgat	H...	Anheuser-Busch

*Machine 1*

*Machine 2*

*Machine 5*

```
SELECT name  
FROM Beers  
WHERE manf = 'Heineken'
```

*Broadcast query*

*In each machine in parallel:*

*Select<sub>manf='Heineken'</sub> (Beers)*

*Project(name)*

*Gather Partial Results*

*Union*

*Return*

# Local and Global Indexing

- What if a machine does not have any data for the query attributes?
- Index structures
  - Given value, return records
  - Several solutions
    - Use local index on each machine
    - Use a machine index for each value
    - Use a combined index in a global index server

manf	RecordIDs
...	...
MillerCoors	34, 35, 87, 129, ...
Duvel Moortgat	5, 298, 943, 994, ...
Heineken	631, 683, 882, ...
...	...

manf	machineIDs
...	...
MillerCoors	10
Duvel Moortgat	3, 4
Heineken	1, 3, 5
...	...

# Pause

# Querying Two Relations

- Often we need to combine two relations for queries

- Find the beers liked by drinkers who frequent The Great American Bar

*Frequents(drinker, bar)*  
*Likes(drinker, beer)*

- In SQL

- SELECT DISTINCT beer
  - FROM Likes L, Frequents F
  - WHERE bar = The Great American Bar' AND
  - F.drinker = L.drinker

# SPJ Queries

*Frequents(drinker, bar)*  
*Likes(drinker, beer)*

*SELECT DISTINCT beer*  
*FROM Likes L, Frequents F*  
*WHERE bar = 'The Great American Bar'*  
*AND F.drinker = L.drinker*

- **Steps**

Selection<sub>bar = 'The Great American Bar'</sub> (Frequents)

Join<sub>F.drinker = L.drinker</sub> (⋈, Likes) *No intermediate storage*

Project<sub>beer</sub>(⋈) *R(drinker, beer)*

Deduplicate(⋈)

Output

# Join in a Distributed Setting

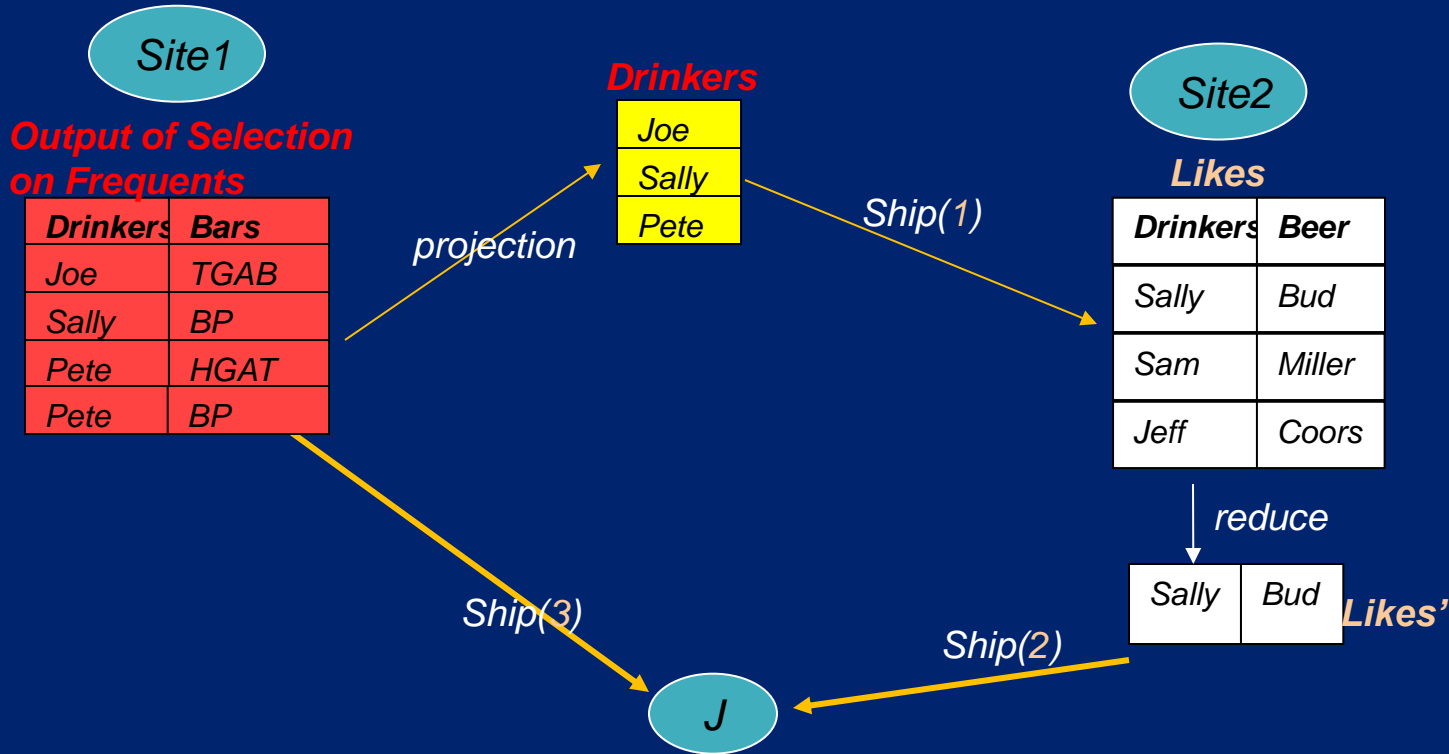
*Frequents(drinker, bar)*

*Likes(drinker, beer)*

- **Semijoin**


- A semijoin from R to S on attribute is used to reduce the data transmission cost
- Computing steps:
  - **Project** R on attribute A and call it  $(R[A])$  – the Drinkers column
  - **Ship** this projection ( a semijoin projection) from the site of R to the site of S
  - **Reduce** S to  $S'$  by eliminating tuples where attribute A are not matching any value in  $R[A]$

# Semijoin s: *Frequents*—*Drinkers* → *Likes*





# Semijoin $s$ : $Frequents \rightarrow Drinkers \rightarrow Likes$



# Pause

# Subqueries


- A slightly complex query
- Find the bars that serve Miller for the same or less price than what TGAB charges for Bud
- We may break it into two queries:
  1. Find the price TGAB charges for Bud
  2. Find the bars that serve Miller at that price

# Subqueries in SQL

```
SELECT bar  
FROM Sells  
WHERE beer = 'Miller' AND
```

```
price <= (SELECT price  
FROM Sells  
WHERE bar = 'TGAB'  
AND beer = 'Bud');
```

*The price at  
which TGAB  
sells Bud*



# Subqueries with IN

- Find the name and manufacturer of each beer that Fred does not like

- Query

```
SELECT *
```

```
FROM Beers
```

```
WHERE name NOT IN
```

```
    ( SELECT beer
```

```
        FROM Likes
```

```
        WHERE drinker = 'Fred');
```

*Beers(name, manf)*  
*Likes(drinker, beer)*

# Correlated Subqueries

- Find the name and price of each beer that is more expensive than the average price of beers sold in the bar

```
SELECT beer, price
FROM Sells s1
WHERE price >
  (SELECT AVG(price)
   FROM Sells s2
   WHERE s1.bar = s2.bar)
```

Bar	Beer	Price
HGAT	Bud	5
BP	Michelob	4
TGAB	Heineken	6
HGAT	Guinness	10

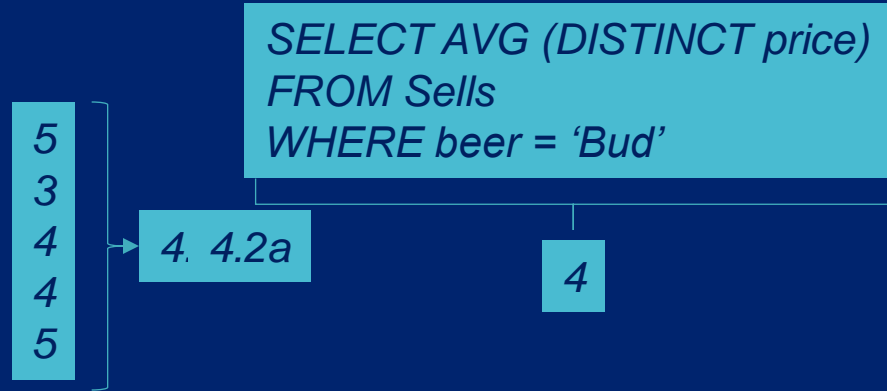
# Aggregate Queries

- **Example**

- Find the average price of Bud:
- `SELECT AVG(price)`
- `FROM Sells`
- `WHERE beer = 'Bud';`

- **Other aggregate functions**

- SUM, MIN, MAX, COUNT, ...



# GROUP BY Queries

- Find for each drinker the average price of Bud at the bars they frequent

```
SELECT drinker, AVG(price)
```

```
FROM Frequents, Sells
```

```
WHERE beer = 'Bud' AND
```

```
    Frequents.bar = Sells.bar
```

```
GROUP BY drinker;
```

Drinker	Bar	Price
Pete	HGAT	5
Pete	BP	4
Joe	TGAB	6
Joe	HGAT	5



Drinker	Price
Pete	4.5
Joe	5.5



# Grouping Aggregates over Partitioned Data

Drinker	Bar	Price
Pete	HGAT	5
Pete	BP	4
Joe	TGAB	6
John	HGAT	5

Drinker	Bar	Price
Pete	HGAT	5
Pete	BP	4
Pete	BO	6
Joe	TGAB	6

Drinker	Price
Pete	5
Joe	6

Drinker	Bar	Price
Pete	BO	6
John	BP	4
Sally	TGAB	6
Sally	HGAT	5

Drinker	Bar	Price
John	HGAT	5
John	BP	4
Sally	TGAB	6
Sally	HGAT	5

Drinker	Price
John	4.5
Sally	5.5